

# Assignment\_4

February 28, 2025

## 1 Assignment 4: Software for neural network training

By: Tim Bakkenes, Elvira Moberg and Agnes Magnusson

```
[22]: import pandas as pd
import numpy as np
import torch
```

## 2 Task 1

Loading the synthetic dataset.

```
[23]: # You may need to edit the path, depending on where you put the files.
data = pd.read_csv('data/a4_synthetic.csv')

X = data.drop(columns='y').to_numpy()
Y = data.y.to_numpy()
```

Training a linear regression model for this synthetic dataset.

```
[24]: np.random.seed(1)

w_init = np.random.normal(size=(2, 1))
b_init = np.random.normal(size=(1, 1))

# We just declare the parameter tensors. Do not use nn.Linear.
w = torch.tensor(w_init, dtype=torch.float, requires_grad=True)
b = torch.tensor(b_init, dtype=torch.float, requires_grad=True)

eta = 1e-2
opt = torch.optim.SGD([w, b], lr=eta)

for i in range(10):

    sum_err = 0

    for row in range(X.shape[0]):
```

```

        x = torch.tensor(X[[row], :], dtype=torch.float) # added dtype=torch.
↪float to ensure same format
        y = torch.tensor(Y[[row]], dtype=torch.float)      # added dtype=torch.
↪float to ensure same format

        # Forward pass.
        y_pred = x @ w + b
        err = (y - y_pred)**2

        # Backward and update.
        opt.zero_grad()
        err.backward()
        opt.step()

        # For statistics.
        sum_err += err.item()

mse = sum_err / X.shape[0]
print(f'Epoch {i+1}: MSE =', mse)

```

```

Epoch 1: MSE = 0.7999662647869263
Epoch 2: MSE = 0.017392394159767264
Epoch 3: MSE = 0.009377418162580966
Epoch 4: MSE = 0.009355327616258364
Epoch 5: MSE = 0.009365440349979508
Epoch 6: MSE = 0.009366988411857164
Epoch 7: MSE = 0.009367207068114567
Epoch 8: MSE = 0.009367238481529512
Epoch 9: MSE = 0.009367244712136654
Epoch 10: MSE = 0.009367244620257224

```

### 3 Task 2

```

[25]: class Tensor:

        # Constructor. Just store the input values.
        def __init__(self, data, requires_grad=False, grad_fn=None):
            self.data = data
            self.shape = data.shape
            self.grad_fn = grad_fn
            self.requires_grad = requires_grad
            self.grad = None

        # So that we can print the object or show it in a notebook cell.
        def __repr__(self):
            dstr = repr(self.data)
            if self.requires_grad:

```

```

        gstr = ', requires_grad=True'
    elif self.grad_fn is not None:
        gstr = f', grad_fn={self.grad_fn}'
    else:
        gstr = ''
    return f'Tensor({dstr}{gstr})'

# Extract one numerical value from this tensor.
def item(self):
    return self.data.item()

# YOUR WORK WILL BE DONE BELOW

# For Task 2:

# Operator +
def __add__(self, right):
    # Call the helper function defined below.
    return addition(self, right)

# Operator -
def __sub__(self, right):
    return subtraction(self, right)

# Operator @
def __matmul__(self, right):
    return matrix_multiplication(self, right)

# Operator **
def __pow__(self, right):
    # NOTE! We are assuming that right is an integer here, not a Tensor!
    if not isinstance(right, int):
        raise Exception('only integers allowed')
    if right < 2:
        raise Exception('power must be >= 2')
    else:
        return power(self, right)

# Backward computations. Will be implemented in Task 4.
def backward(self, grad_output=None):
    # We first check if this tensor has a grad_fn: that is, one of the
    # nodes that you defined in Task 3.
    if self.grad_fn is not None:
        # If grad_fn is defined, we have computed this tensor using some
        ↪ operation.
        if grad_output is None:

```

```

        # This is the starting point of the backward computation.
        # This will typically be the tensor storing the output of
        # the loss function, on which we have called .backward()
        # in the training loop.

        #initialize grad output to 1s
        grad_output = np.ones_like(self.data)

        self.grad_fn.backward(grad_output)
    else:
        # This is an intermediate node in the computational graph.

        # This corresponds to any intermediate computation, such as
        # a hidden layer.

        #call the backward function recursively
        self.grad_fn.backward(grad_output)
    else:
        # If grad_fn is not defined, this is an endpoint in the
        computational
        # graph: learnable model parameters or input data.
        if self.requires_grad:
            # This tensor *requires* a gradient to be computed. This will
            # typically be a tensor that holds learnable parameters.

            #save grad_output in self.grad
            self.grad = grad_output

        else:
            # This tensor *does not require* a gradient to be computed.

            # This
            # will typically be a tensor holding input data.

            #terminate the recursion
            return

# A small utility where we simply create a Tensor object. We use this to
# mimic torch.tensor.
def tensor(data, requires_grad=False):
    return Tensor(data, requires_grad)

# We define helper functions to implement the various arithmetic operations.

# This function takes two tensors as input, and returns a new tensor holding
# the result of an element-wise addition on the two input tensors.
def addition(left, right):

```

```

    new_data = left.data + right.data
    grad_fn = AdditionNode(left, right)
    return Tensor(new_data, grad_fn=grad_fn)

def subtraction(left, right):
    new_data = left.data - right.data
    grad_fn = SubtractionNode(left, right)
    return Tensor(new_data, grad_fn=grad_fn)

def matrix_multiplication(left, right):
    new_data = left.data @ right.data
    grad_fn = MatrixMultiplicationNode(left, right)
    return Tensor(new_data, grad_fn=grad_fn)

def power(left, right):
    new_data = left.data**right
    grad_fn = PowerNode(left, right)
    return Tensor(new_data, grad_fn=grad_fn)

def tanh(x):
    new_data = np.tanh(x.data)
    grad_fn = TanhNode(x)
    return Tensor(new_data, grad_fn=grad_fn)

def sigmoid(x):
    # Ensure x.data is a numpy array
    data = np.asarray(x.data)
    new_data = 1 / (1 + np.exp(-data))
    grad_fn = SigmoidNode(x)
    return Tensor(new_data, grad_fn=grad_fn)

def BCELoss(y_pred, y_true):
    #using sigmoid to convert y_pred from logits to probabilities
    y_prob = sigmoid(y_pred)
    new_data = -y_true.data*np.log(y_prob.data) - (1-y_true.data)*np.
    ↪log(1-y_prob.data)
    grad_fn = BCELossNode(y_prob, y_true)
    return Tensor(new_data, grad_fn=grad_fn)

```

Some sanity checks.

```

[27]: # Two tensors holding row vectors.
x1 = tensor(np.array([[2.0, 3.0]]))
x2 = tensor(np.array([[1.0, 4.0]]))
# A tensors holding a column vector.
w = tensor(np.array([[ -1.0], [ 1.2]]))

```

```

# Test the arithmetic operations.
test_plus = x1 + x2
test_minus = x1 - x2
test_power = x2 ** 2
test_matmul = x1 @ w

print(f'Test of addition: {x1.data} + {x2.data} = {test_plus.data}')
print(f'Test of subtraction: {x1.data} - {x2.data} = {test_minus.data}')
print(f'Test of power: {x2.data} ** 2 = {test_power.data}')
print(f'Test of matrix multiplication: {x1.data} @ {w.data} = {test_matmul.
↳data}')

# Check that the results are as expected. Will crash if there is a
↳miscalculation.
assert(np.allclose(test_plus.data, np.array([[3.0, 7.0]])))
assert(np.allclose(test_minus.data, np.array([[1.0, -1.0]])))
assert(np.allclose(test_power.data, np.array([[1.0, 16.0]])))
assert(np.allclose(test_matmul.data, np.array([[1.6]])))

```

```

Test of addition: [[2. 3.]] + [[1. 4.]] = [[3. 7.]]
Test of subtraction: [[2. 3.]] - [[1. 4.]] = [[ 1. -1.]]
Test of power: [[1. 4.]] ** 2 = [[ 1. 16.]]
Test of matrix multiplication: [[2. 3.]] @ [[-1. ]
[ 1.2]] = [[1.6]]

```

Create some tensors and make sure that you can compute the arithmetic operations that you used in the linear regression examples above.

## 4 Tasks 3 and 4

For each node, the derivative of the corresponding function is calculated and passed backwards, enabling backpropagation.

```

[28]: class Node:
    def __init__(self):
        pass

    def backward(self, grad_output):
        raise NotImplementedError('Unimplemented')

    def __repr__(self):
        return str(type(self))

class AdditionNode(Node):
    def __init__(self, left, right):
        self.left = left
        self.right = right

```

```

    def backward(self, grad_output):
        self.left.backward(grad_output)
        self.right.backward(grad_output)

class SubtractionNode(Node):
    def __init__(self, left, right):
        self.left = left
        self.right = right

    def backward(self, grad_output):
        self.left.backward(grad_output)
        self.right.backward(-grad_output)

class MatrixMultiplicationNode(Node):
    def __init__(self, left, right):
        self.left = left
        self.right = right

    def backward(self, grad_output):
        self.left.backward(grad_output @ self.right.data.T)
        self.right.backward(self.left.data.T @ grad_output)

class PowerNode(Node):
    def __init__(self, left, right):
        self.left = left
        self.right = right

    def backward(self, grad_output):
        self.left.backward(grad_output*self.right*(self.left.data**(self.
↪right-1)))

class TanhNode(Node):
    def __init__(self, x):
        self.x = x

    def backward(self, grad_output):
        self.x.backward(grad_output*(1-(np.tanh(self.x.data)**2)))

class BCELossNode(Node):
    def __init__(self, y_prob, y_true):
        self.y_prob = y_prob
        self.y_true = y_true

    def backward(self, grad_output):
        derivative = (self.y_prob.data - self.y_true.data)/(self.y_prob.
↪data*(1-self.y_prob.data))

```

```

        self.y_prob.backward(grad_output*derivative)

class SigmoidNode(Node):
    def __init__(self, x):
        self.x = x

    def backward(self, grad_output):
        self.x.backward(grad_output*(self.x.data*(1-self.x.data)))

```

Sanity check for Task 3.

```

[29]: x = tensor(np.array([[2.0, 3.0]]))
w1 = tensor(np.array([[1.0, 4.0]]), requires_grad=True)
w2 = tensor(np.array([[3.0, -1.0]]), requires_grad=True)

test_graph = x + w1 + w2

print('Computational graph top node after x + w1 + w2:', test_graph.grad_fn)

assert(isinstance(test_graph.grad_fn, AdditionNode))
assert(test_graph.grad_fn.right is w2)
assert(test_graph.grad_fn.left.grad_fn.left is x)
assert(test_graph.grad_fn.left.grad_fn.right is w1)

```

Computational graph top node after x + w1 + w2: <class '\_\_main\_\_.AdditionNode'>

Sanity check for Task 4.

```

[30]: x = tensor(np.array([[2.0, 3.0]]))
w = tensor(np.array([[ -1.0], [1.2]]), requires_grad=True)
y = tensor(np.array([[0.2]]))

# We could as well write simply loss = (x @ w - y)**2
# We break it down into steps here if you need to debug.

model_out = x @ w
diff = model_out - y
loss = diff ** 2

loss.backward()

print('Gradient of loss w.r.t. w =\n', w.grad)

assert(np.allclose(w.grad, np.array([[5.6], [8.4]])))
assert(x.grad is None)
assert(y.grad is None)

```

Gradient of loss w.r.t. w =  
[[5.6]



[8.4]]

An equivalent cell using PyTorch code. Your implementation should give the same result for `w.grad`.

```
[31]: pt_x = torch.tensor(np.array([[2.0, 3.0]]))
      pt_w = torch.tensor(np.array([[ -1.0], [ 1.2]]), requires_grad=True)
      pt_y = torch.tensor(np.array([[0.2]]))

      pt_model_out = pt_x @ pt_w
      pt_model_out.retain_grad() # Keep the gradient of intermediate nodes for
      ↪ debugging.

      pt_diff = pt_model_out - pt_y
      pt_diff.retain_grad()

      pt_loss = pt_diff ** 2
      pt_loss.retain_grad()

      pt_loss.backward()
      pt_w.grad
```

```
[31]: tensor([[5.6000],
              [8.4000]], dtype=torch.float64)
```

## 5 Task 5

```
[32]: class Optimizer:
      def __init__(self, params):
          self.params = params

      def zero_grad(self):
          for param in self.params:
              param.grad = np.zeros_like(param.data)

      def step(self):
          raise NotImplementedError('Unimplemented')

      class SGD(Optimizer):
          def __init__(self, params, lr):
              super().__init__(params)
              self.lr = lr

          def step(self):
              for param in self.params:
                  if param.grad is not None:
                      if param.grad.shape[0] != param.data.shape[0]:
```

```

        #sum gradients for bias, then apply gradient descent
        grad_b = np.sum(param.grad, axis=0, keepdims=True)
        param.data -= self.lr * grad_b
    else:
        #gradient descent on each weight
        param.data -= self.lr * param.grad

```

Testing if the same results are achieved with the code form task 1

```

[33]: np.random.seed(1)

w_init = np.random.normal(size=(2, 1))
b_init = np.random.normal(size=(1, 1))

# We just declare the parameter tensors. Do not use nn.Linear.
w = tensor(w_init, requires_grad=True)
b = tensor(b_init, requires_grad=True)

eta = 1e-2

opt = SGD([w, b], lr=eta)

for i in range(10):

    sum_err = 0

    for row in range(X.shape[0]):
        x = tensor(X[[row], :])
        y = tensor(Y[[row]])

        # Forward pass.
        y_pred = x @ w + b
        err = (y - y_pred)**2
        #err = err.mean()

        # Backward and update.
        opt.zero_grad()
        err.backward()
        opt.step()

        # For statistics.
        sum_err += err.item()

    mse = sum_err / X.shape[0]
    print(f'Epoch {i+1}: MSE =', mse)

```

Epoch 1: MSE = 0.7999661130823178

Epoch 2: MSE = 0.017392390107906875

```
Epoch 3: MSE = 0.009377418010839892
Epoch 4: MSE = 0.009355326971438456
Epoch 5: MSE = 0.009365440968904256
Epoch 6: MSE = 0.009366989180952533
Epoch 7: MSE = 0.009367207398577986
Epoch 8: MSE = 0.009367238983974489
Epoch 9: MSE = 0.009367243704122532
Epoch 10: MSE = 0.009367244427185763
```

## 6 Task 6

```
[34]: from sklearn.preprocessing import scale
      from sklearn.model_selection import train_test_split

      # You may need to edit the path, depending on where you put the files.
      a4data = pd.read_csv('data/raisins.csv')

      X = scale(a4data.drop(columns='Class'))
      Y = 1.0*(a4data.Class == 'Besni').to_numpy()

      np.random.seed(0)
      shuffle = np.random.permutation(len(Y))
      X = X[shuffle]
      Y = Y[shuffle]

      Xtrain, Xtest, Ytrain, Ytest = train_test_split(X, Y, random_state=0,
      ↪test_size=0.2)
```

Initialize the model structure with one hidden layer using tanh as activation function

```
[35]: class RaisinCLF:
      def __init__(self, input_dim, hidden_dim, output_dim):
          #initialize weights with small random values, and biases with zero.
          ↪Requires_grad=true to enable backpropagation
          self.w1 = tensor(np.random.randn(input_dim, hidden_dim) * 0.01,
          ↪requires_grad=True)
          self.b1 = tensor(np.zeros((1, hidden_dim)), requires_grad=True)
          self.w2 = tensor(np.random.randn(hidden_dim, output_dim) * 0.01,
          ↪requires_grad=True)
          self.b2 = tensor(np.zeros((1, output_dim)), requires_grad=True)

      def forward(self, x):
          """forward pass through network"""

          #computing hidden layer with linear transformation and tanh activation
          hidden = tanh(addition(matrix_multiplication(x, self.w1), self.b1))
```

```

        #computing output layer with linear transformation
        output = addition(matrix_multiplication(hidden, self.w2), self.b2)
        return output

    def get_params(self):
        return [self.w1, self.b1, self.w2, self.b2]

```

```

[36]: from sklearn.model_selection import train_test_split
      from sklearn.metrics import accuracy_score

def train_classifier(model, X, Y):
    # hyperparameters
    lr = 1e-6
    n_epochs = 100
    val_size = 0.2

    Xtrain, Xval, Ytrain, Yval = train_test_split(X, Y, test_size=val_size,
↪random_state=0)

    optimizer = SGD(model.get_params(), lr=lr)

    #stores training and validation accuracy
    history = []

    #convert data to tensor
    Xtrain_t = tensor(Xtrain)
    Ytrain_t = tensor(Ytrain)

    max_acc = 0

    for epoch in range(n_epochs):

        loss_sum = 0

        outputs = model.forward(Xtrain_t)

        loss = BCELoss(outputs, Ytrain_t)

        #initilize gradient to 0
        optimizer.zero_grad()

        #computing gradient
        loss.backward()

        #apply gradient descent
        optimizer.step()

```

```

        #compute average loss
        loss_sum += loss.data.mean()

        train_pred = predict(model, Xtrain).flatten()
        val_pred = predict(model, Xval).flatten()

        train_acc = accuracy_score(Ytrain.flatten(), train_pred)
        val_acc = accuracy_score(Yval.flatten(), val_pred)

        max_acc = max(max_acc, val_acc)

        history.append((train_acc, val_acc))

        if (epoch+1) % 5 == 0:
            print(f'Epoch {epoch+1}: loss = {loss_sum:.4f}, train acc = {train_acc:.4f}, val acc = {val_acc:.4f}')

            print(f'Max validation accuracy achieved: {max_acc:.4f}')
            return history

def predict(model, x):
    Xt = tensor(x)
    scores = sigmoid(model.forward(Xt).data)
    y_guess = (scores.data > 0.5).astype(int)
    return y_guess

```

```

[37]: input_dim = 7
      hidden_dim = 128
      output_dim = 1

      model = RaisinCLF(input_dim=input_dim, hidden_dim=hidden_dim,
                        output_dim=output_dim)

      history = train_classifier(model, Xtrain, Ytrain.reshape(-1, 1))

      y_guess = predict(model, Xtest)

      accuracy = accuracy_score(Ytest, y_guess)
      print(f'Test accuracy: {accuracy:.4f}')

```

```

Epoch 5: loss = 0.6921, train acc = 0.8594, val acc = 0.8194
Epoch 10: loss = 0.6921, train acc = 0.8594, val acc = 0.8194
Epoch 15: loss = 0.6921, train acc = 0.8576, val acc = 0.8125
Epoch 20: loss = 0.6921, train acc = 0.8576, val acc = 0.8194
Epoch 25: loss = 0.6921, train acc = 0.8594, val acc = 0.8194
Epoch 30: loss = 0.6921, train acc = 0.8594, val acc = 0.8194
Epoch 35: loss = 0.6921, train acc = 0.8594, val acc = 0.8194

```

Epoch 40: loss = 0.6921, train acc = 0.8594, val acc = 0.8194  
Epoch 45: loss = 0.6921, train acc = 0.8576, val acc = 0.8194  
Epoch 50: loss = 0.6921, train acc = 0.8576, val acc = 0.8194  
Epoch 55: loss = 0.6921, train acc = 0.8576, val acc = 0.8194  
Epoch 60: loss = 0.6921, train acc = 0.8576, val acc = 0.8125  
Epoch 65: loss = 0.6921, train acc = 0.8576, val acc = 0.8125  
Epoch 70: loss = 0.6921, train acc = 0.8576, val acc = 0.8125  
Epoch 75: loss = 0.6921, train acc = 0.8559, val acc = 0.8125  
Epoch 80: loss = 0.6921, train acc = 0.8576, val acc = 0.8125  
Epoch 85: loss = 0.6921, train acc = 0.8576, val acc = 0.8194  
Epoch 90: loss = 0.6921, train acc = 0.8594, val acc = 0.8194  
Epoch 95: loss = 0.6921, train acc = 0.8611, val acc = 0.8194  
Epoch 100: loss = 0.6921, train acc = 0.8611, val acc = 0.8264  
Max validation accuracy achieved: 0.8264  
Test accuracy: 0.8333  
Epoch 100: loss = 0.6921, train acc = 0.8611, val acc = 0.8264  
Max validation accuracy achieved: 0.8264  
Test accuracy: 0.8333